

What Is This Project About?

You're building a **web crawler and search engine system** — like a mini version of Google — but using **distributed cloud computing**. Instead of running on one computer, it runs on **many virtual machines in the cloud** (like AWS, GCP, Azure) that work together to:

- **Crawl websites** (download pages)
 - **Index content** (organize it so it can be searched)
 - **Handle failure** (if a part of the system breaks, it keeps working)
-

What Needs to Be Implemented?

1. Client Interface

- A simple **web interface or CLI** (Command Line Interface) to:
 - Start a web crawl (by giving URLs)
 - Set options (like how deep to crawl, limit to domains)
 - Search the index for keywords
 - Monitor the progress
-

2. Master Node (Controller)

This is the brain of the system:

- **Splits URLs into tasks** and sends them to workers
 - **Schedules tasks** for crawler/indexer VMs
 - **Watches workers' health** (are they running or failed?)
 - **Reassigns tasks** if something goes wrong
 - **Manages the index** creation
-

3. Crawler Nodes (Worker VMs)

These download pages:

- **Fetch pages** from assigned URLs
 - **Parse the page:** extract text, links, and metadata
 - **Follow rules** like robots.txt
 - **Send data** to indexer nodes
 - **Report status** to master
-

4. Indexer Nodes

These build the searchable index:

- **Take content** from crawlers
 - **Create index** (e.g., inverted index = word → list of pages)
 - **Store index** in **cloud storage**
 - **Answer search queries**
-

5. Distributed Task Queue

- A **cloud-based message queue** (like AWS SQS or GCP Pub/Sub)
 - Used for:
 - Sending tasks from master → crawlers/indexers
 - Receiving results/status from workers
-

6. Cloud Storage

- Stores everything:
 - **Seed URLs**
 - **Downloaded HTML pages**
 - **Processed content**
 - **Index**
- Should be **durable and safe** (doesn't lose data)

Extra Features to Implement

Scalability

- You should be able to **add more VMs** when needed to crawl/index faster.

Fault Tolerance

Make sure the system keeps working even if some parts fail:

- **Crawler failure:**
 - Master detects it (e.g., via heartbeat signals)
 - Reassigns the URLs to other crawlers
- **Indexer failure:**
 - Duplicate/store parts of the index
 - Rebuild lost data if needed
- **Reliable queues and storage:**
 - Use built-in reliability of cloud services

User Stories (What Users Can Do)

- Start a crawl with a list of URLs
- Set crawl settings (like how deep to go)
- Keep crawling even if some machines fail
- Search for words in indexed data
- Track system progress (crawl status)




Summary of What You'll Build

Component	What it Does
Client Interface	Start crawl, search content, monitor progress
Master Node	Controls the whole system, assigns work, handles faults
Crawler Nodes	Download and parse web pages
Indexer Nodes	Build and store the search index
Task Queue	Handles task distribution and communication
Cloud Storage	Stores data and index safely

Let's say the website is:

 <https://news.site.com>

This website might have many pages, such as:

-  <https://news.site.com/home>
-  <https://news.site.com/sports>
-  <https://news.site.com/tech/article123>
-  <https://news.site.com/world/article456>

Each of those is a **separate web page** under the same website.

Here's how it works step by step:

- 1. You provide seed URLs**
 - For example:

<https://example.com>

<https://news.site.com>
- 2. The system crawls those URLs**
 - It downloads those pages
 - Extracts links from them
 - Follows those links (up to a certain depth or within domain limits)

3. The content from crawled pages is indexed

- The system builds a search index (like a map of words to pages)

4. Later, you can search with keywords

- After crawling is done, you can search for terms like:

"climate change"

"machine learning"

- And the system will return the relevant pages it crawled.